

COMS W4701: Artificial Intelligence

Lecture 2: Search Problems

Tony Dear, Ph.D.

Department of Computer Science
School of Engineering and Applied Sciences

Today

- Search problem formulation
- State space graphs and search trees
- Uninformed search: DFS, BFS, UCS
- Informed search: Greedy, A*
- Search heuristics: Admissibility, design
- Applet for self-studying: <http://www.aispace.org/search/index.shtml>

Problem-Solving Agents


- Goal-based agent whose goal is a **state**, or description of the current task environment
- States contain all relevant information, can be treated as black boxes
- Environment properties: **fully observable, single-agent, deterministic, static, discrete**
 - Percepts are trivial, since we see entire environment
 - Action results always known, go from one state to another state
- Agent wants to find an *action sequence* that will result in a *state sequence* to a goal
- This is the agent's solution to a **search problem**

Search Problems

- **State space S :** Set of descriptions of the agent and environment
- **Actions:** (Finite) set of available actions in a state
 - Ex: $Actions(s_1) = \{a_1, a_2, a_3\}$
- **Transition model:** (Deterministic) mapping from (state, action) to a new state
 - Ex: $Result(s_1, a_1) = s_2$
- **Action costs:** Numerical cost for a (state, action, new state) transition
 - Ex: $Cost(s_1, a_1, s_2) = 10$
- **Goal test (for goal states)**
 - Ex: $IsGoal(s_1) = \text{False}$, $IsGoal(s_2) = \text{True}$

Example: Grid World Path Finding

- **State space:** Current coordinates of the rabbit
 - $S = \{(x, y) \mid 0 \leq x \leq 3, 0 \leq y \leq 4\}$
- **Actions:** $Actions((x, y)) = \{\text{Up, Down, Left, Right}\}$
- **Costs:** $Cost(s, a, s') = 1, \forall s, a, s'$
- **Transition model:** $Result((x, y), \text{Up}) = (x, y + 1), Result((x, y), \text{Down}) = \dots$
 - Should also account for walls and boundaries, e.g. $Result((0,0), \text{Left}) = (0,0)$
- **Goal test:** $In((3,3))?$



Search Problem Example: n -puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State


- **State:** Locations of all tiles and blank
- **Action:** 4 possible directions for the blank tile
- **Action cost:** Each step taken costs 1
- **Goal test:** Is current state equal to goal state?

More Search Problems

- Route-finding (e.g., vehicle navigation), robot navigation in the real world
- Touring problems (traveling salesperson)
- Layout and assembly sequencing problems
- Mathematical puzzles and proofs: Infinitely large state spaces!
- Knuth's conjecture (1964): Starting with the number 4, use a combination of factorial, floor, and sqrt operations to reach any other desired integer
- States: All nonnegative integers


State Space Graphs

- **State space graph:** A mathematical representation of a search problem
 - *Vertices* are states; *edges* are actions
 - Each state occurs only once!
- *Paths* are sequences of actions/states
- A *solution* is a path from initial to goal state
- We can rarely build this full graph in memory—it can be very large or infinite




Search Trees

- Need a systematic way of performing search over a state space graph
- **Search tree:** Nodes are states, edges are actions; root is initial state



- Unlike state space graph, states can occur more than once
- Each node corresponds to a *unique* path from initial state

General Search Ideas



- From current node, **expand** and consider all possible actions
- Generate successor **nodes** for each resultant state according to transition function
 - Each node should track its corresponding state, parent, prior action, and total cost so far
- Successors are added to a **frontier** of possible next nodes to expand
- Frontier forms a boundary between explored and unexplored parts of tree

Node Expansion

- We still have some unanswered questions re: node expansion...
- How to select the next node from the frontier?
- **Best-first search:** Implement frontier as a priority queue; each node is assigned a priority according to an *evaluation function* $f(n)$ (lowest priority node popped first)
 - **Uninformed search:** $f(n)$ has no knowledge about how close a state is to goal
- Suppose we expand a node and a child node has already been expanded...
- We may want to consider it again if this new occurrence is through a cheaper path!
- Idea: Keep track of all *reached* nodes in a lookup table

Best-First Search

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
  node  $\leftarrow$  NODE(STATE=problem.INITIAL)
  frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
  reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    for each child in EXPAND(problem, node) do
      s  $\leftarrow$  child.STATE
      if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
        reached[s]  $\leftarrow$  child
        add child to frontier
  return failure
```


```
function EXPAND(problem, node) yields nodes
  s  $\leftarrow$  node.STATE
  for each action in problem.ACTIONS(s) do
    s'  $\leftarrow$  problem.RESULT(s, action)
    cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
    yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

Depth-First Search

- Idea: Expand the *deepest* node in the frontier (relative to start node)
- Node priority can be implemented as the *negative of depth*
- Alternatively, the frontier can be implemented as a stack (LIFO)
- No consideration of true costs! DFS only “cares” about depth information
- Possible optimizations from best-first search base implementation:
 - *Early goal test*: Check if node is goal upon insertion into frontier (instead of removal)
 - Do not use a reached table, possibly exploring a node more than once

DFS Properties

- **Time complexity:** How many nodes to explore in the worst case? $O(b^m)$
- **Space complexity:** How many frontier nodes to keep in memory? $O(bm)$
 - Assumes no need for “reached” table
- **Completeness:** Not if state space is infinite
- **Optimality:** No, only returns first solution



- b is the *branching factor*
- m is the *maximum depth*
- Total nodes: $O(1 + b + b^2 + \dots + b^m)$

Breadth-First Search

- Idea: Expand the *shallowest* node in the frontier (relative to start node)
- Node priority is exactly equal to the node *depth*
- Alternatively, the frontier can be implemented as a FIFO queue
- No consideration of true costs! BFS only “cares” about depth information
- BFS may be optimal if true costs are equivalent to depths
- Possible optimizations from best-first search base implementation:
 - *Early goal test*: Check if node is goal upon insertion into frontier (instead of removal)
 - Reached table can just be a set of states (no need to track costs)

BFS Properties

- **Time complexity:** How many nodes to explore in the worst case?


$$O(b^d)$$

- **Space complexity:** How many frontier nodes to keep in memory?

$$O(b^d)$$

- **Completeness:** If solution exists, yes!


- **Optimality:** Only if costs are uniform



- d is depth of the shallowest solution
- May be significantly smaller than m
- Max frontier size is $O(b^d)$



Improving DFS and BFS

- **Depth-limited DFS:** Prevent DFS from going past a set depth l
- Time complexity $O(b^l)$, space complexity $O(bl)$
- Best if we know *diameter* of state space in advance
- **Iterative-deepening:** Iteratively do depth-limited search with increasing l : try $l = 0$, then $l = 1, \dots$
- Ends when l reaches d (depth of shallowest solution)
- Time complexity $O(b^d)$, space complexity $O(bd)$
- Why is wasted effort in upper levels of search tree not a concern?





Uniform-Cost Search (Dijkstra)

- Idea: Expand node that least increases total cost
- Evaluation function $f(n)$ is the *cumulative path cost*
- For UCS to be correct, we must use a reached table and conduct a late goal test



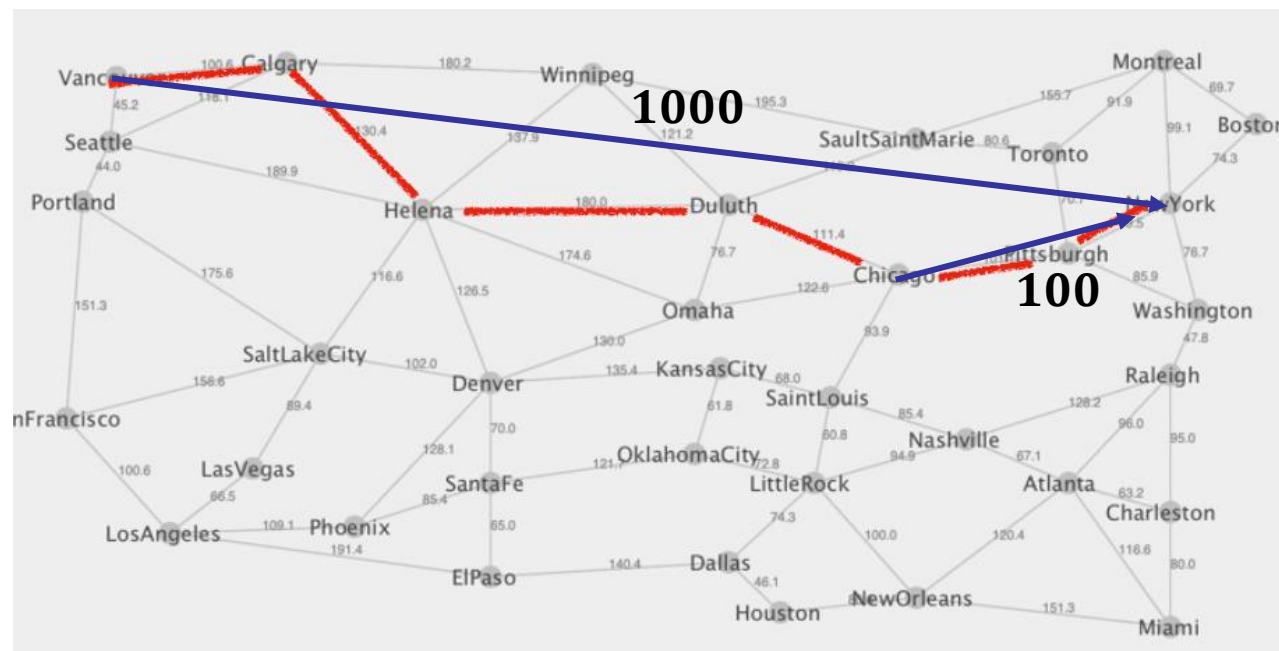
UCS Properties

- Let C^* be the cost of optimal solution
- Let ϵ be lower bound on all possible costs
- $1 + C^*/\epsilon$ is the max depth to traverse before finding optimal solution
- Time and space complexity: $O(b^{1+C^*/\epsilon})$
- UCS is both **complete** and **optimal**



Informed (Heuristic) Search


- Oftentimes we have additional, *domain-specific heuristics* that tell us how close a state is to a goal
- **Heuristic function $h(n)$:** Estimated cost of cheapest path from state at node n to a goal state
- Often come from *relaxed problems*, precomputed *subproblem* solutions, or learning from experience



Example: Euclidean or Manhattan distance on a map


Greedy Best-First Search

- Idea: Expand node that *appears* closest to goal according to heuristic function
- Evaluation function is just the heuristic function! $f(n) = h(n)$
- As with DFS and BFS, there is no consideration of *true* costs



A* Search

- Idea: From a given node, estimate the *best* path that *continues* to the goal
- $f(n) = g(n) + h(n)$: Sum of path cost to n and estimated cost from n to goal
- Benefits of both UCS and greedy best-first search







Node $f(n) = g(n) + h(n)$


S	$6 = 0 + 6$
a	$6 = 1 + 5$
b	$8 = 2 + 6$
e	$10 = 9 + 1$
f	$7 = 4 + 3$
d	$9 = 8 + 1$
c	$10 = 3 + 7$
G	$9 = 9 + 0$

A* vs UCS vs BFS

- BFS expands search tree by increasing depth
- UCS expands acc. to increasing g -cost
- Contours are “circular” around start state, if normalized by path costs
- A* expands acc. to increasing $g + h$ cost
- If heuristic is good, expanded states should show preference toward goal



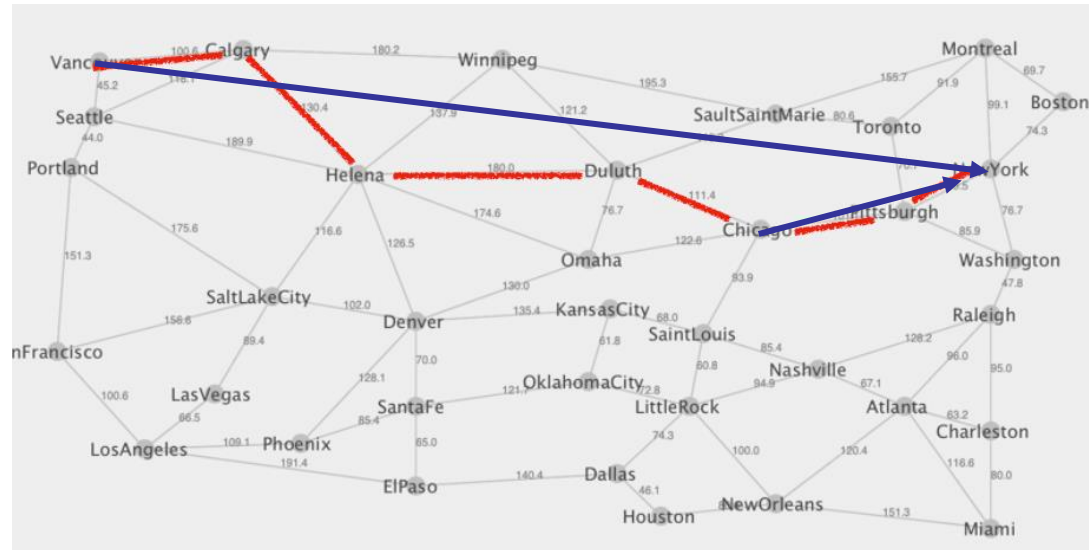
When is A* Optimal?



- What is the problem here?
- Heuristic along optimal path overestimated the true cost!
- Good heuristics should be optimistic—never overestimate true costs

Admissible Heuristics

- A heuristic h is **admissible** if $0 \leq h(n) \leq h^*(n)$ where $h^*(n)$ is true cost from n to goal
- Most heuristics derived from relaxed problems are admissible
- Same state space graph, but with added edges
- With fewer constraints or restrictions, problems are easier to solve
- Example: Euclidean distances



Misplaced Tiles Heuristic

- $h(n)$ = number of misplaced tiles, not including blank

	1	2
3	4	5
6	7	8

$h(\text{state}) = 0$

1	4	2
	5	8
3	6	7

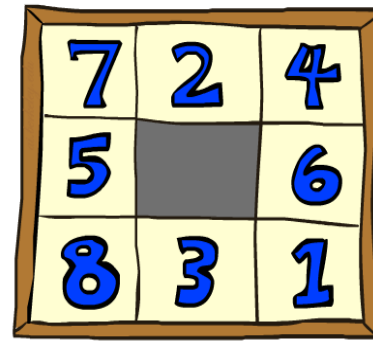
$h(\text{state}) = 7$

- Relaxed problem: Any tile can be correctly replaced with just one move
- Admissible because misplaced tiles will always require *at least* one move

Manhattan Distance Heuristic

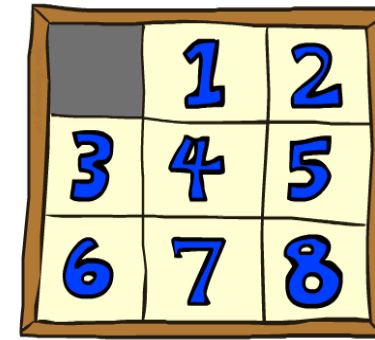
- $h(n)$ = sum of Manhattan distances between current tile positions and goal positions

$$h(start) = 18$$



7	2	4
5		6
8	3	1

Start State



	1	2
3	4	5
6	7	8

Goal State


- Relaxed problem: Multiple tiles can simultaneously occupy same space
- Admissible because misplaced tiles will always require *at least* number of moves equal to Manhattan distance

Heuristic Domination

- For any node n , Manhattan distance heuristic $h_2(n) >$ misplaced tiles heuristic $h_1(n)$
- h_2 **dominates** h_1 if $h_2(n) \geq h_1(n)$ for all n
- A* search using h_2 will be more efficient and never expand more nodes than h_1
- h_2 reflects true costs more accurately
- Suppose we have collection of admissible heuristics h_1, h_2, \dots, h_m
- The composite heuristic $h(n) = \max\{h_1(n), \dots, h_m(n)\}$ is admissible and dominates all other heuristics!

Completeness and Optimality of A*

- A* is complete (same reason as UCS, BFS)
- If heuristic function is admissible, A* is also optimal!
- Suppose A* returns a suboptimal solution
- Then there exists some unexpanded node n on optimal path
- Since n was not expanded, $f(n) > C^*$ (optimal cost)
- By definition, $f(n) = g(n) + h(n)$: backward cost + heuristic
- Since h is admissible, $h(n) \leq h^*(n)$ (true cost-to-go), and so $f(n) \leq g(n) + h^*(n)$
- But $C^* = g(n) + h^*(n)$, meaning that $f(n) \leq C^*$. Contradiction!



Other A* Points

- Performance of A* depends entirely on choice of heuristic function
- Assuming an admissible heuristic, worst case time and space complexity is the same as UCS (all heuristics equal to 0)
- Otherwise, we can only improve from UCS
- Heuristics and the use of domain knowledge make A* distinctly more “intelligent” than its uninformed counterparts
- While still hard theoretically, problems become easier in practice

Satisficing Solutions*

- Like BFS or UCS, A* may suffer computationally intractable memory requirements
- Idea: Trade off admissibility for more accurate heuristics to reduce computation
- Return **satisficing solutions**—suboptimal, but “good enough”
- **Weighted A* search:** $f(n) = g(n) + \alpha h(n)$
- We can choose to place higher weight α on the heuristic
- Generalizes A* ($\alpha = 1$), UCS ($\alpha = 0$), and greedy best-first ($\alpha = \infty$)
- Suboptimality: If optimal solution has cost C^* , weighted A* solution may cost up to αC^*

Memory-Bounded Search*

- We can also consider A^* variants that are more memory-efficient
- **Beam search:** Limit frontier size by discarding worst nodes past a given limit
- Alternatively, discard nodes with scores much smaller than best one
- **Iterative-deepening A^* (IDA*):** Repeatedly run A^* with increasing depth limit
- Nodes with higher f -cost than limit are treated as leaves
- Increment depth limit by smallest f -cost of “leaves” from previous iteration
- IDA* worst case: Each node has different f -cost, num iterations equal to num states

Summary

- Objective of search problems is to find action/state sequence to reach a goal state
- Represented by state space graphs; search algorithms follow a tree structure
- Uninformed search: No usage of information indicating closeness to goal
- Examples: Depth-first, breadth-first, depth-limited, iterative deepening, uniform-cost
- Generally suffer from lack of completeness or intractable memory usage
- Informed search: Domain-specific heuristics guide search toward goal
- Greedy best-first and A* search use a heuristic function to evaluate frontier nodes
- Optimal if heuristics are admissible: good, optimistic estimates of true costs