

COMS W4701: Artificial Intelligence

Lecture 3: Constraint Satisfaction Problems

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today

- Constraint satisfaction problems
- Backtracking search
- Ordering heuristics
- Inference and constraint propagation
- Local consistency

States with Structure


- Task environments so far: Fully observable, discrete, deterministic, static
- Transitions and heuristics are problem-specific
- States are *atomic* black boxes
- **Planning** solution is an action **sequence**, or state space **path**
- What if our states have a common **factored** representation?
- We can use *general-purpose* heuristics
- Solution of an **assignment** problem is the **goal** itself, not a path

Constraint Satisfaction Problems

- Special structured search problems with 3 components
 - Variables $X = \{X_1, \dots, X_n\}$
 - Domains $D = \{D_1, \dots, D_n\}$
 - Constraints $C = \{C_1, \dots, C_m\}$
- Goal test: A **complete, consistent** assignment of values to each variable X_i from respective domain D_i s.t. all constraints C are satisfied
- We may have incomplete or inconsistent assignments along the way


Example: Map Coloring

- Goal: Color a map so that no adjacent territories have the same color
- Variables: $X = \{WA, NT, Q, NSW, V, SA, T\}$
- Domains: $D_i = \{\text{red, green, blue}\}$
- Constraints: Implicit vs explicit representation
 - $C = \{WA \neq NT, WA \neq SA, NT \neq SA, NT \neq Q, SA \neq Q, Q \neq NSW, NSW \neq V\}$
 - $C = \{(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), \dots\}, \dots\}$



Constraint Graphs

- Visualization of a CSP in a graph
- Nodes: Variables and domains
- Edges: Presence of binary constraints
- Certain graph algorithms can help give insight into the CSP
- E.g., a k -connected component indicates a k -nary constraint
- CSPs that are actually tree structures can be solved without backtracking




Types of Constraints

- So far we've seen **binary** constraints relating two variables
- **Unary** constraints involve a single variable, equivalent to *domain reduction*
 - Ex: $SA \neq \text{green}$ (implicit), $SA \in \{\text{green}, \text{red}\}$ (explicit)
- Higher-order (**global**) constraints relate arbitrary number of variables
 - Ex: *Alldiff* requiring all variables to have different values
- If domains are finite, always possible to rewrite global constraints as binary ones using *auxiliary variables*

Example: n -Queens


- Place n queens on $n \times n$ board s.t. none share a row, column, or diagonal
- Variables: X_{ij} , $1 \leq i \leq n$, $1 \leq j \leq n$ (grid spaces)
- Domains: $D_{ij} = \{0,1\}$ (queen or no queen)



- Constraints:
$$\forall i \sum_j X_{ij} = 1 \quad \forall i, j \sum_k X_{i+k, j+k} \leq 1$$
$$\forall j \sum_i X_{ij} = 1 \quad \forall i, j \sum_k X_{i+k, j-k} \leq 1$$

Example: n -Queens

- Alternatively, states can just represent a row of the board (vs grid space)
- Variables: $X_i, 1 \leq i \leq n$ (row of the board)
- Domains: $D_i = \{1, \dots, n\}$ (column containing queen)
- Constraints: $\forall i, j, X_i \neq X_j$
 $\forall i, j, X_j - X_i \neq |j - i|$



Example: Cryptarithmic

- Variables: $X = \{T, W, O, F, U, R, \dots\}$
- Also carry-overs! C_1, C_2, C_3

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

- Domains: $D_i = \{0, \dots, 9\}$

- Constraints

- $Alldiff(X_i)$
 - $$O + O = 10C_1 + R$$

$$W + W + C_1 = 10C_2 + U$$

$$T + T + C_2 = 10F + O$$

$$C_3 = F$$

- 7 possible solutions!

$$734 + 734 = 1468$$

$$765 + 765 = 1530$$

$$836 + 836 = 1672$$

$$846 + 846 = 1692$$

$$867 + 867 = 1734$$

$$928 + 928 = 1856$$

$$938 + 938 = 1876$$

Example: Sudoku

- Variables: One for each open square
- Domains: $\{1, \dots, 9\}$
- Constraints:
 - *Alldiff* for each column
 - *Alldiff* for each row
 - *Alldiff* for each 3×3 square

					8			4
	8	4		1	6			
			5			1		
1		3	8			9		
6		8				4		3
		2			9	5		1
		7			2			
			7	8		2	6	
2			3					

Solving CSPs with Search


- **States:** Partial assignments (initial state is no assignment)
- **Actions:** Assign value to an unassigned variable from its domain
- **Goal test:** Complete, consistent assignment
- No explicit costs, so maybe something like DFS/BFS?
- Problem: Naïve implementation -> lots of repeated states!
- **Branching factor:** Number of unassigned variables \times size of domain
 - Branching factor at root: $n \times d$ for n variables with domain size d
 - Branching factor at 2nd level is $(n - 1)d$, then $(n - 2)d$ at 3rd level, ...
 - Total number of leaves / possible goal states: $O(n! \times d^n)$

Backtracking Search

- Idea: CSPs are *commutative*, order of variable assignment doesn't matter
 - (WA = red, NT = green) is the same as (NT = green, WA = red)
- Each search tree tier only needs to correspond to a single variable
 - Branching factor is d at every level, number of leaves is $O(d^n)$
- Running DFS allows us to make one variable assignment at a time and *backtrack* by undoing inconsistent assignments to try alternatives
 - Don't need a frontier—just keep track of what domain values are available

Variable Selection


- Heuristic for picking next variable: Minimum remaining values (**MRV**)
 - Also called “most constrained variable” (MCV) or “fail-first”
 - Choose the variable with the fewest legal left values in its domain



- What if we have multiple MRV variables (e.g., initial assignment)?
 - *Degree* heuristic: Pick variable that appears in the most constraints
 - Reduce branching factor on future choices

Domain Value Assignment

- Heuristic for assigning a value: Least constraining value (**LCV**)
 - Choose a value that imposes fewest constraints on future assignments
 - May require some computation
- Why least, not most?
 - We don't have to use all values, just find an assignment that works
 - Try to look for the most likely ones earlier
 - Keep more options open for other variables

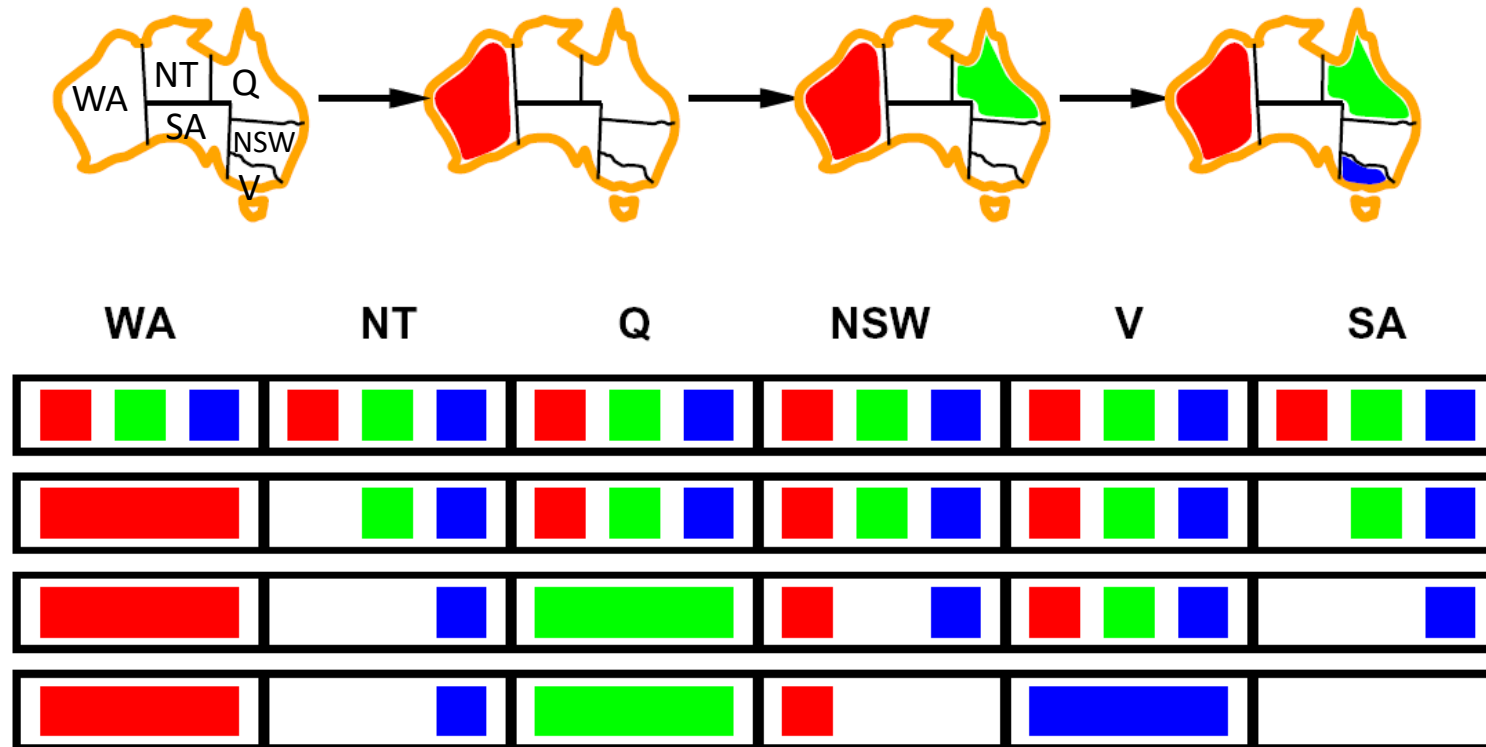


Improving Backtracking: Inference

- Both MRV and LCV heuristics assume some sort of “forward checking”
- We can detect inevitable failure earlier by checking constraints as we go
- Idea: Delete inconsistent domain values as we make assignments
- If any domain is left empty, we should not proceed
- Domain reduction also better informs MRV on selecting variables
- Inference can also inform how to use the LCV heuristic

Forward Checking

- **Forward checking:** After assigning a variable, check other variables related to it by a constraint and eliminate inconsistent domain values
- If any domain becomes empty, current assignment is denoted a failure



Local Consistency

- Forward checking is an example of **constraint propagation**
- Can be interleaved with search, or done as a preprocessing step prior to search
- Result is **local consistency**: remaining domain values do not violate any constraints

- **Node consistency**: All unary constraints are satisfied
 - Can always be done prior to starting any search process


- **Arc consistency**: No binary constraint is violated by any variable's domain values
 - X_i is arc-consistent with respect to X_j if every value in D_i can be paired with a value in D_j such that no constraint between X_i and X_j is violated.

Arc Consistency

- X_i being arc consistent with X_j does not imply that X_j is arc consistent with X_i !
- Example: $Y = X^2$, both with domains $\{0,1, \dots, 9\}$
 - To make X arc consistent with Y : Reduce domain of X to $\{0,1,2,3\}$
 - To make Y arc consistent with X : Reduce domain of Y to $\{0,1,4,9\}$
- Entire CSP is arc consistent iff both domains are reduced
- Another issue: Arc consistency may not be preserved if a domain is reduced
- If X_i is arc consistent with X_j and X_j is changed (e.g. due to forward checking), X_i *may no longer be arc consistent with X_j*



AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check





AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check





AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check





AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check




AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check



AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check



Path and k -Consistency

- At least one arc inconsistent \Rightarrow no solution
 - BUT arc consistency not sufficient to guarantee a solution!
 - Idea: Look at more than two variables at a time
-
- The diagram illustrates a constraint network with three nodes. The top node is connected to two bottom nodes. Each node contains two colored squares (red and blue). The top node has a red square on the left and a blue square on the right. The bottom-left node has a red square on the left and a blue square on the right. The bottom-right node has a red square on the left and a blue square on the right. The nodes are connected by lines, forming a tree structure.
- $\{X_i, X_j\}$ is **path consistent** with X_m if for every arc consistent assignment to $\{X_i, X_j\}$, there is an assignment to X_m consistent with X_i and X_j
 - **k -consistency**: For every consistent assignment to a set of $k - 1$ variables, a consistent value can be assigned to the k th variable
 - A **strongly n -consistent** CSP ($n, n - 1, \dots, 1$ consistent) is guaranteed a solution, but determining n -consistency is exponentially hard (NP-complete)!

Special Constraints


- Consistency for certain constraint types can be checked quickly
- *Alldiff* with m variables and n unique values: immediately declare no solution if $m > n$
- Suppose we have domains with lower and upper bounds
- **Bounds propagation:** If constraints are equality or inequality (e.g., **resource constraints**), we can use them to tighten the bounds and make them consistent
- Ex: X_1, X_2, X_3 , all with domains $[1,5]$
 - Constraint $\sum X_i = 13$: Reduce domains to $[3,5]$
 - Constraint $\sum X_i \leq 5$: Reduce domains to $[1,3]$

Introducing New Constraints

- We can also improve CSP solvers by *introducing* new constraints
- Recall: We backtrack whenever we see an inconsistent solution, either in the present (leaf of the search tree) or in the future (via constraint propagation)
- **Constraint learning:** Record the current assignment as a constraint
- A CSP may have multiple solutions due to *value symmetry*
- Ex: In map coloring, there are $d!$ solutions by permutation
- Idea: Introduce **symmetry-breaking constraints** (e.g., $NT < SA < WA$) to reduce the number of solutions and shrink the search tree

Local Search


- Methods so far: Build up solution incrementally, check constraints
- **Local search:** Start with an arbitrary complete assignment, *modify* it until consistent
- No frontier to maintain, no backtracking!



- **Min-conflicts:** Reassign a variable to a value that *minimizes conflicts*

```
function MIN-CONFLICTS(csp, max_steps) returns a solution or failure  
  inputs: csp, a constraint satisfaction problem  
           max_steps, the number of steps allowed before giving up  
  
  current  $\leftarrow$  an initial complete assignment for csp  
  for i = 1 to max_steps do  
    if current is a solution for csp then return current  
    var  $\leftarrow$  a randomly chosen conflicted variable from csp.VARIABLES  
    value  $\leftarrow$  the value v for var that minimizes CONFLICTS(csp, var, v, current)  
    set var = value in current  
  return failure
```


Example: 4-Queens



- Can be very effective in practice, roughly independent of problem size
- E.g., can solve the 1M-queens problem in an average of 50 iterations
- Also useful for solving problems online as they change
- For example, repair existing solutions as new constraints come and go

Local Search and Optimization*

- Local search generally useful when we care about goal more than path
- E.g., optimization problems with **objective function**, no specific goal test
- Pros: Very little memory, can work well in large or infinite state spaces




Hill-Climbing Methods*

- Many local search algorithms fall into category of *hill-climbing*
- No search tree, no knowledge beyond immediate neighbors
- Greedy search method, can get stuck at local maxima if unlucky


- Variations of simply ascending the hill:
 - **Stochastic**: Pick random moves from time to time
 - **First-choice**: Pick a random move that leads to a better successor
 - **Random-restart**: Try different initial states if stuck

Other Methods*

- **Simulated annealing:** Combine hill climbing with random walks to get unstuck from local maxima
- **Local beam search:** Maintain parallel searches but share information among threads
- **Genetic algorithms:** Generate successor states by combining two parent states



<https://eat.zesty.com/post/how-simulated-annealing-can-improve-your-lunch>



Summary

- CSPs are assignment search problems with very specific structure
- Unlike general DFS, backtracking search can be very effective
- General-purpose heuristics: MRV, LCV
- Inference can further improve search performance
- Maintain local consistency through constraint propagation, take advantage of or add new common constraints
- Local search: Alternative search method with few memory requirements