

COMS W4701: Artificial Intelligence

Lecture 4: Adversarial Search and Games

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Where are we now?

- Planning and identification (assignment) problems
- Applications: Pathing, motion planning, scheduling, ...
- Solution: General-purpose search algorithms
 - NP-hard (exponential) in the worst / naïve case
 - Different strategies to make search more efficient
- We have only considered **single-agent** problems

Today

- Adversarial search problems
- Minimax algorithm
- Alpha-beta pruning, evaluation functions, move ordering
- Stochastic games

AI and Games

- 1950s and 60s: First checkers programs, some from reinforcement learning
- 1994: *Chinook* declared computer champion in checkers against Tinsley
- 2007: Checkers is solved (completely predictable)

- 1997: *Deep Blue* defeats chess world champion Kasparov
- 2017, 2019: *AlphaZero*, *Leela Chess Zero* defeat *Stockfish*, AI chess champ

- 2016: *AlphaGo* defeats go world champion Lee Sedol
- 2018-2020: *AlphaZero*, *MuZero* achieve state-of-the-art in chess, go, shogi

Adversarial Search

- Several approaches for competitive multi-agent environments
- Economics: Study aggregate system, no consideration of individual agents
- Can consider other agents as part of a *nondeterministic* environment
- Does not consider goals and motivations of adversaries
- Incorporate other agents into unified **adversarial search**
- Will often have to operate real-time, accept suboptimal solutions

Two-player Zero-sum Games

- Simple adversarial search problem: Two agents, turn-taking, deterministic, perfect information (fully observable), zero-sum (p1 wins = p2 loses and vice-versa)
- States (positions): Current state of the game
- Actions: Set of legal moves in a state s
- Transition model: Mapping from (state, action) to a new state
- **Terminal test**: Is the current state a *terminal state* (is the game over)?
- **Utility function**: A player's "score" or payoff at a terminal state
- Since games are zero-sum, we should have $Utility(s, P1) = -Utility(s, P2)$
 - Ex: Tic-tac-toe. X has three in a row, $Utility(s, X) = 1, Utility(s, O) = -1$

Game Trees



MAX (X)



MIN (O)



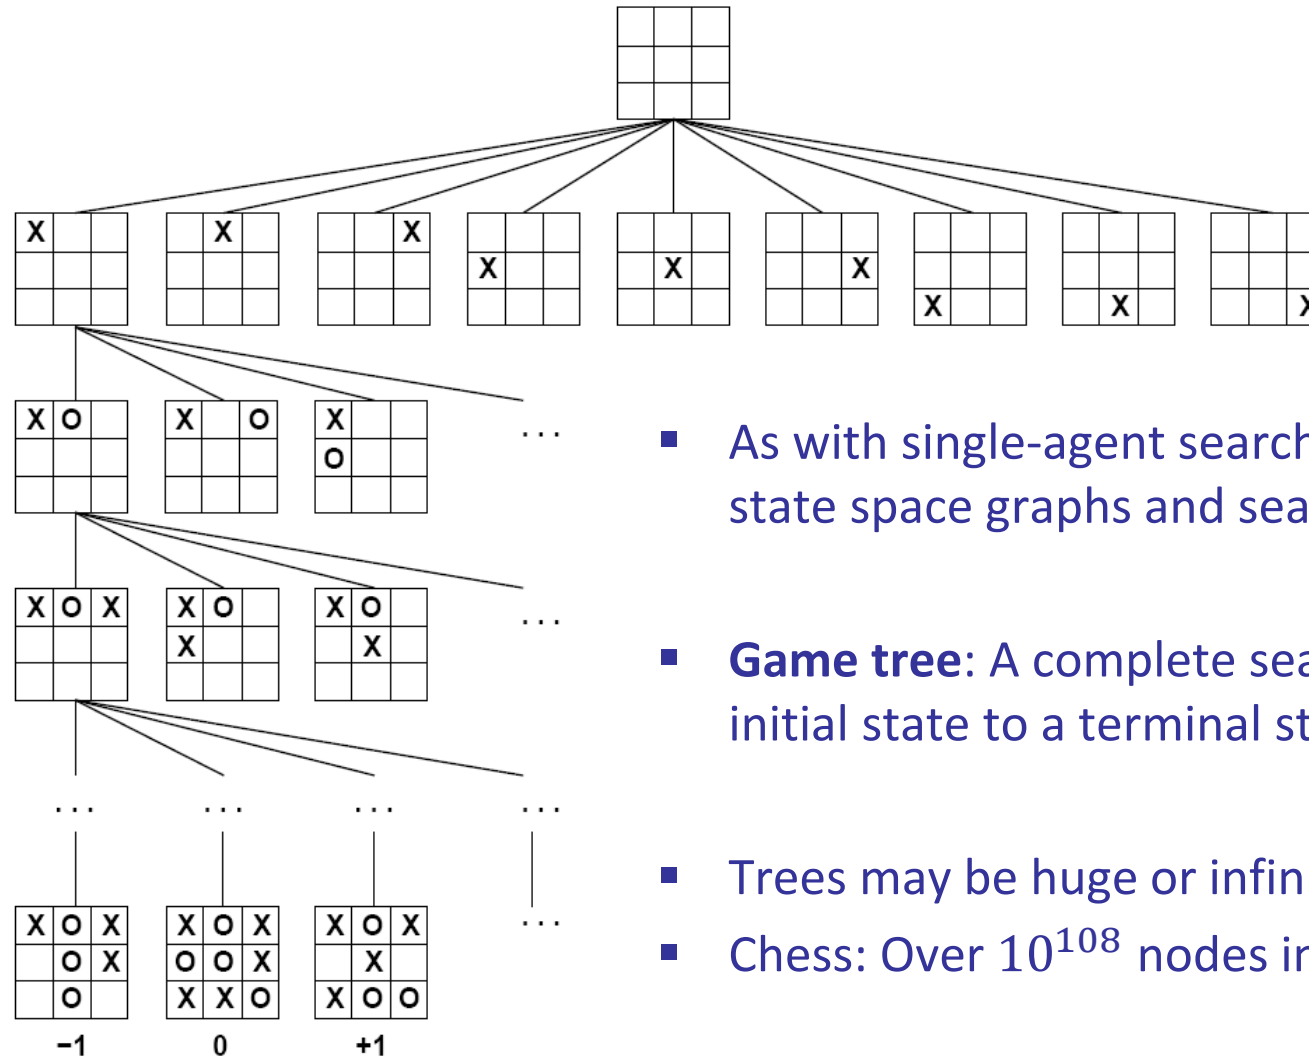
MAX (X)



MIN (O)

TERMINAL

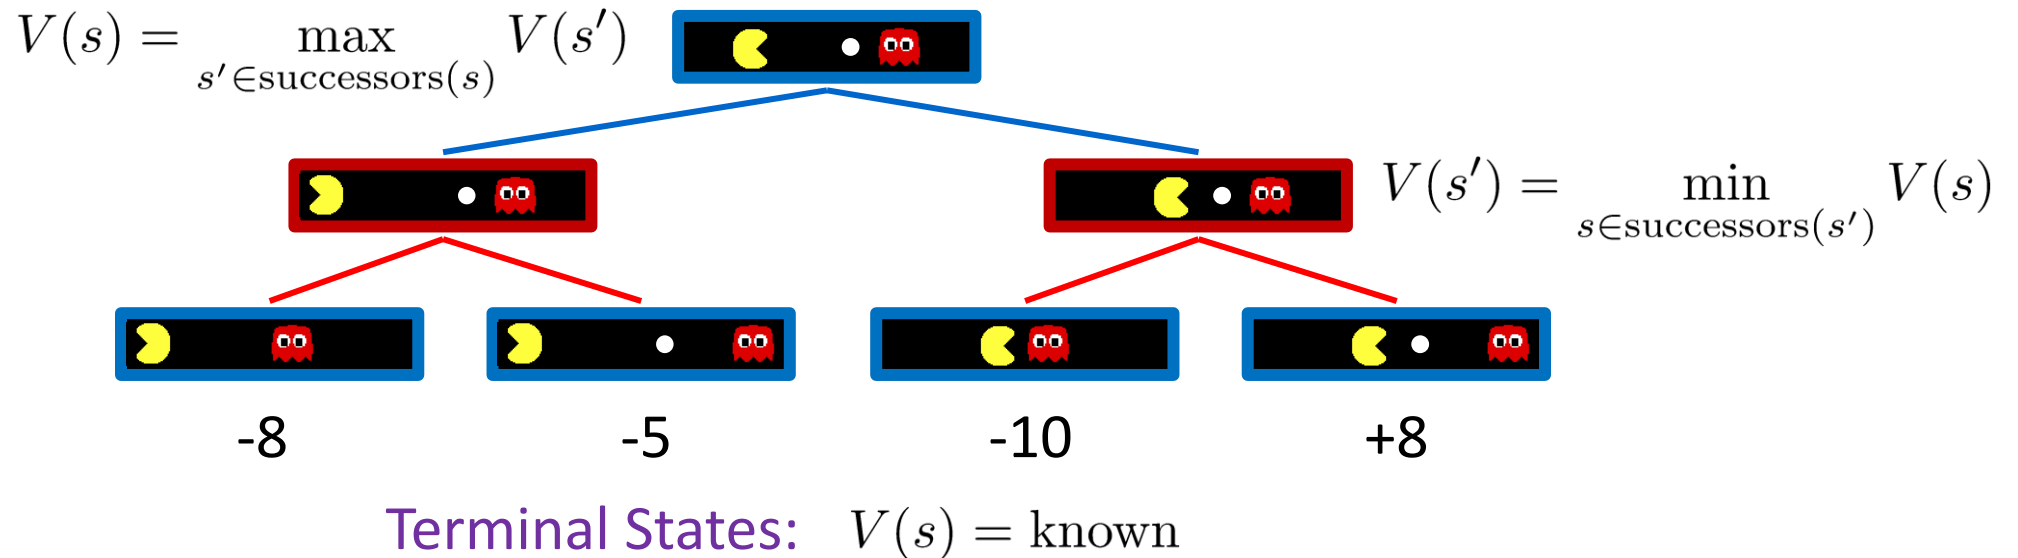
Utility



- As with single-agent search, we can define state space graphs and search trees
- **Game tree:** A complete search tree from an initial state to a terminal state
- Trees may be huge or infinite!
- Chess: Over 10^{108} nodes in the tree

State Utilities

- Both players want to maximize their own utilities
- Equivalently, we can just use MAX's utilities, which player MIN wants to minimize
- Each player should plan its move based on expectation of opponent
- If non-terminal states also have utilities, decision-making would be easy




Minimax Values

- Value of each state depends on each player's move
- If we know that each player plays optimally, the entire game tree's values are known!
- Minimax value:** Utility of a state assuming *both* players play optimally until the end of the game

$\text{MINIMAX}(s) =$


$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax values:
computed recursively



Terminal values:
part of the game

Minimax Example



Minimax Search Algorithm

function MINIMAX-SEARCH(*game*, *state*) **returns** an action

$\text{player} \leftarrow \text{game}.\text{TO-MOVE}(\text{state})$

$\text{value}, \text{move} \leftarrow \text{MAX-VALUE}(\text{game}, \text{state})$ Assuming root is MAX

return *move*

function MAX-VALUE(*game*, *state*) **returns** a (*utility*, *move*) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, *player*), null

$v \leftarrow -\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

$v2, a2 \leftarrow \text{MIN-VALUE}(\text{game}, \text{game}.\text{RESULT}(\text{state}, a))$ MAX calls MIN

if $v2 > v$ **then**

$v, \text{move} \leftarrow v2, a$

return *v*, *move*

function MIN-VALUE(*game*, *state*) **returns** a (*utility*, *move*) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, *player*), null

$v \leftarrow +\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

$v2, a2 \leftarrow \text{MAX-VALUE}(\text{game}, \text{game}.\text{RESULT}(\text{state}, a))$ MIN calls MAX


if $v2 < v$ **then**

$v, \text{move} \leftarrow v2, a$

return *v*, *move*

Multiplayer Games

- We can generalize minimax to non-zero-sum or multiplayer games
- Keep track of *vector* of utilities at each node
- Each player maximizes own utility
- Can produce cooperation, alliances




Improving Minimax

- Minimax executes a DFS-style search of the game tree
- Time complexity $O(b^d)$, space complexity $O(bd)$
- Optimal if both players play perfectly, complete if game eventually ends

- Example: Chess has $b \approx 35$, $d \approx 70$
- Completely infeasible to solve entirely!


- Do we always need to expand entire tree?
- Improvements: Pruning, move ordering, cutting off search

Pruning




Pruning

- Best value for node **E** so far is n (from one of its children)
- If $n \leq m'$, MAX would never choose to go from **C** to **E**
 - $\max(m', \min(n, \dots), \dots) \geq m'$
- Similarly, if $n \leq m$ from node **B**, MAX would never choose the sequence of actions going from **A** to **E**
- In both cases, any remaining children of **E** should be pruned away



Alpha-Beta Pruning

- General idea: Keep track of highest (α) and lowest (β) values seen so far by MAX and MIN nodes, respectively
- Skip remaining children (prune) if:
 - MAX sees value higher than β
 - MIN sees value lower than α
- Know that root node would never choose path to current node



Alpha-Beta Search

function ALPHA-BETA-SEARCH(*game*, *state*) **returns** an action

$\text{player} \leftarrow \text{game.TO-MOVE}(\text{state})$

$\text{value}, \text{move} \leftarrow \text{MAX-VALUE}(\text{game}, \text{state}, -\infty, +\infty)$

return *move*

Assuming root is MAX

function MAX-VALUE(*game*, *state*, α , β) **returns** a (*utility*, *move*) pair

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow -\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MIN-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a), \alpha, \beta)$

if $v2 > v$ **then**

$v, \text{move} \leftarrow v2, a$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if $v \geq \beta$ **then return** *v*, *move*

return *v*, *move*

MAX updates α , compares against β

function MIN-VALUE(*game*, *state*, α , β) **returns** a (*utility*, *move*) pair

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow +\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MAX-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a), \alpha, \beta)$

if $v2 < v$ **then**

$v, \text{move} \leftarrow v2, a$


$\beta \leftarrow \text{MIN}(\beta, v)$

if $v \leq \alpha$ **then return** *v*, *move*

return *v*, *move*

MIN updates β , compares against α


Alpha-Beta Example



- Applet: <http://homepage.ufp.pt/jtorres/ensino/ia/alfabeta.html>

Alpha-Beta Properties

- Pruning does not change true minimax value of the root
- Intermediate (children) node values might be wrong!!
- If pruning, game tree values cannot be reused
- Will have to rerun minimax after each move
- In practice, we can store just the states/values that we know to be correct in a *transposition table* in case they come up again
 - Especially effective when there are multiple paths to a state




Move Ordering

- Good *move ordering* improves effectiveness of pruning
- Alpha-beta with random ordering is roughly $\sim O(b^{0.75d})$
- “Perfect ordering” gets us to $O(b^{0.5d})$, doubling solvable depth
- Usually requires domain knowledge
 - Simple chess ordering function: captures, threats, forward moves, backward moves
 - Try moves first that were good in past moves
- If using depth-limited search, iterative deepening can also inform move ordering
- Evaluations at depth 1 inform ordering at depth 2, those inform ordering at depth 3, and so on

Imperfect Decisions

- Problem: Most game trees still too big
- α - β can help but still need to find terminal nodes
- Heuristic: Turn non-terminal nodes into terminals!
- **Evaluation function** returns an *estimate* of this “terminal” state’s utility
- **Cutoff test** decides when to do this
- No more guarantee of optimality




Evaluation Functions

- Evaluation functions are *estimates* of a state's utility
- Agent's performance depends strongly on eval function quality
 - Evaluation of terminals should be the same as that of true utility function
 - Evaluation of non-terminals should have *some* correlation with winning
 - Computation must be efficient
- One common eval function: weighted linear sum of game **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$




Example: Chess



$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- Features may be derived from expert knowledge of common *categories* of states
- E.g., one feature for each type of piece, attack formations, king safety positions, etc.
- Weights correspond to *material values* of each feature
- Linear weighting assumes features are independent of each other

Example: Go



$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- Example features in Go:
 - Num white pieces – num black pieces
 - Buildup of potential “ladders”
 - Territorial “spread” of pieces

Depth Limits


- Now that we have eval function, when to cut off search and apply it?
- Simple approach: Use a fixed depth limit, or use iterative deepening
- Transposition table especially effective with latter approach
- Problem: Cutting off search at volatile positions can lead to loss of information
- **Horizon effect:** Agent may favor moves that push danger “over the horizon”, appears to have been mitigated but actually just delayed
- If possible, only cut off search at **quiescent** (quiet) positions
- **Quiescence search:** Extend search at volatile positions until we find quieter ones
- “Adaptive” search depth based on domain knowledge

Stochastic Games


- Many games contain *stochasticity*
- Opponents playing suboptimally (e.g., inexperienced) or randomly
- Explicit random elements (e.g., dice rolling)
- Instead of worst-case scenarios, we consider **expected values**
- Chance nodes have **expectiminimax values**

EXPECTIMINIMAX(s) =


$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$



Example: Backgammon




Expectiminimax Example




Expectiminimax Utilities

- Decisions in stochastic games are sensitive to choice of eval function
- Minimax: Decisions unchanged as long as relative ordering of node values is the same
- Expectiminimax: No guarantee that decisions stay the same if values change
- To fix, must ensure that eval function is a *positive linear transformation* of true utilities



Expectiminimax Pruning

- Pruning is more difficult than in deterministic games
- Cannot eliminate possibilities if we cannot predict results of stochastic transitions
- Pruning *can* occur if we have explicit bounds on utility values overall
- If we can accept imperfect decisions, we can perform forward pruning to examine fewer branches




All utilities ≤ 10

Monte Carlo Tree Search*

- Minimax and its variants are mostly **type A** strategies: search wide but shallow
- For games like Go with large branching factors, **type B** strategies like MCTS work better: search deep but narrow
- Idea: Run many simulations from current game state to a terminal
- No heuristic evaluations; we record end game results in all those simulations
- Results are tabulated and averaged; moves are chosen to maximize chances of winning
- Simulations may be informed by **selection policy** and **playout policy**


Selection Policy

- We grow a tree starting with current game state at the root
- Keep track of how many plays and wins we have simulated from each node
- A predetermined **selection policy** guides us down the tree to a leaf
- Goal: Obtain more simulation results for promising nodes (*exploitation*) but also for lesser known nodes (*exploration*)




Playout Policy

- From selected leaf, add a single child node to the tree
- Simulate a complete game from the child state
- Encountered game states are *not* recorded!
- How do we perform a simulation in our head?
- Follow a predetermined **playout policy**
- Generally biases moves toward good or clever ones
- May incorporate game-specific heuristics
- May be obtained from deep learning (e.g. AlphaGo)



Backpropagation

- Once a simulation is done, the result is returned up the tree
- All nodes along path get updated
- Process iterates between growing and updating search tree
- When done, return the most simulated move
- Why not return move with highest win rate?
- We must also be wary of uncertainty!



MCTS vs Alpha-Beta

- MCTS simulations are linear in game depth
- For a game with branching factor of 32 and depth of 100, alpha-beta search down to 12 ply deep is equivalent to 10^7 MCTS simulations
- MCTS tends to do better when branching factor is high
- Also less sensitive than α - β to inaccurate eval functions
- Also good for brand new games with no predefined eval functions at all!
- Stochastic nature of MCTS: no guarantee of exploring all good moves

Summary

- Adversarial search problems consider the goals of multiple agents
- Two-player zero-sum games can be solved via minimax search
- Improvements: Alpha-beta pruning, move ordering, transposition tables, evaluation functions and depth limits for imperfect decisions
- Can be extended to stochastic games with expectiminimax search
- Generally much harder, need to consider many more search possibilities