

COMS W4701: Artificial Intelligence

Lecture 6: Dynamic Programming

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

Today

- Bellman optimality equations
- Dynamic programming for MDPs
- Value iteration
- Policy iteration


Markov Decision Processes

MDPs: Stochastic, sequential decision problems

- Set of states S , set of actions A
- Transitions $T(s, a, s') = \Pr(s'|s, a)$
- Rewards $R(s, a, s')$, discount γ
- We see state-action-reward $(s, a, r, s, a, r, \dots)$ sequences

Can derive the following functions

- **Policy** $\pi: S \rightarrow A$, assignment of action to each state
- **Value** $V^\pi: S \rightarrow \mathbb{R}$, expected state utilities if following π




Recursive Relationship

$$V^\pi(s) = E \left[\sum_{t=0} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \right]$$

- This is a function on every state in the state space
- For a given state s , we can alternatively write $V^\pi(s)$ as a *recursive* function of *successor state values*

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$


- $V^\pi(s)$ is a *weighted average* of (immediate reward plus discounted successor state values)



Solving for Values

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- Suppose we know the model (T, R) , discount γ , and a fixed policy π
- Then the above is a system of $|S|$ *linear* equations in the $|S|$ unknowns $V^\pi(s)$
- Linear solvers: $\sim O(|S|^3)$ time, can find $V^\pi(s)$



Example: Mini-Gridworld

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- States: A, B, C ; actions: left, right; $\gamma = 0.5$
- Policy: $\pi(s) = \text{left } \forall s$
- Rewards: $R(s, a, A) = +3, R(s, a, B) = -2, R(s, a, C) = +1$
- Transitions: $\text{Pr}(\text{intended direction}) = 0.8, \text{Pr}(\text{opposite direction}) = 0.2$
- V^π can be found by solving 3 linear equations:

+3	-2	+1
A	B	C

$$V^\pi(A) = 0.8(3 + 0.5V^\pi(A)) + 0.2(-2 + 0.5V^\pi(B))$$

$$V^\pi(B) = 0.8(3 + 0.5V^\pi(A)) + 0.2(1 + 0.5V^\pi(C))$$

$$V^\pi(C) = 0.8(-2 + 0.5V^\pi(B)) + 0.2(1 + 0.5V^\pi(C))$$

Bellman Optimality Equations

- Our goal is to find an **optimal policy** or **optimal value function**

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

$$\pi^* = \operatorname{argmax}_\pi V^\pi$$


$$V^* = \max_\pi V^\pi$$



$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Bellman optimality equations are nonlinear!**




Value Function to Policy

- We don't know (yet) how to solve for V^* from scratch
- We *do* know how to find V^* given π^* (solve linear system)

- Bellman equation tells us how to find π^* given V^*

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Everything on the RHS is known!
- Solving for complete policy takes $O(|S|^2|A|)$ time



Example: Mini-Gridworld

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

+3	-2	+1
----	----	----

- States: A, B, C ; actions: left, right
- Transitions and rewards same as before

- Given $V^*(A) = 4.06, V^*(B) = 4.36, V^*(C) = 1.39$

- Find $\pi^*(B)$:


$$\gamma = 0.5 \quad \pi^*(B) = \operatorname{argmax} \begin{cases} 0.8(3 + 0.5V^*(A)) + 0.2(1 + 0.5V^*(C)) & \text{Left} \\ 0.8(1 + 0.5V^*(C)) + 0.2(3 + 0.5V^*(A)) & \text{Right} \end{cases}$$

State-Action Values

- We can also tidy up the Bellman equations by defining **state-action values** (Q-values) of **Q-states**
- Interpretation: Agent has committed to an action, but transition has not yet resolved

$$\begin{aligned} Q(s, a) &= \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \\ &= \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma \max_{a'} Q(s', a')] \end{aligned}$$

$$V^*(s) = \max_a Q(s, a) \quad \pi^*(s) = \operatorname{argmax}_a Q(s, a)$$



Solving the Bellman Equations

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- How to solve for V^* ?
- We have $|S|$ nonlinear equations (because of max)!
- Dynamic programming: *Iterate* on **time-limited values** V_i

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- Bellman equation rewritten as Bellman *update rule*

Value Iteration

Idea: Repeatedly applying *Bellman update* to approximations of value function brings it closer to optimal (true) values V^*

- Initialize: $V_0(s) \leftarrow 0$ for all states s
- **Loop** from $i = 0$:
 - **For** each state $s \in S$:

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- **Until** $\max_s |V_{i+1}(s) - V_i(s)| < \textit{small threshold}$

Example: Mini-Gridworld

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$




- States: A, B, C ; actions: L, R
- Rewards received when entering state
- Transitions: $\Pr(\text{intended direction}) = 0.8, \Pr(\text{opposite direction}) = 0.2$
- Initialize: $(V_0(A), V_0(B), V_0(C)) = (0, 0, 0)$ $\gamma = 0.5$

+3	-2	+1
----	----	----


$$\begin{pmatrix} V_1(A) \\ V_1(B) \\ V_1(C) \end{pmatrix} = \begin{pmatrix} \max[0.8(3 + 0.5(0)) + 0.2(-2 + 0.5(0)), 0.8(-2 + 0.5(0)) + 0.2(3 + 0.5(0))] \\ \max[(0.8(3 + 0.5(0)) + 0.2(1 + 0.5(0))), 0.8(1 + 0.5(0)) + 0.2(3 + 0.5(0))] \\ \max[(0.8(-2 + 0.5(0)) + 0.2(1 + 0.5(0))), 0.8(1 + 0.5(0)) + 0.2(-2 + 0.5(0))] \end{pmatrix} = \begin{pmatrix} 2 \\ 2.6 \\ 0.4 \end{pmatrix}$$

- V_2, V_3, \dots until convergence

Example: Race Car

			
V_2	3.5	2.5	0
V_1	2	1	0
V_0	0	0	0

Assume no discount: $\gamma = 1$



$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

$$V_2(\text{Cool}) = \max(1[1 + 1(2)], 0.5[2 + 1(2)] + 0.5[2 + 1(1)])$$

$$= \max(3, 3.5)$$

$$V_2(\text{Warm}) = \max(0.5[1 + 1(2)] + 0.5[1 + 1(1)], 1[-10 + 1(0)])$$

$$= \max(2.5, -10)$$

Iterative Policy Evaluation

- Iterative update idea can also be used to solve for values of a fixed policy!
- Alternative to solving linear system; no max since actions are fixed

- Initialize $V_0(s)$ for all states s

- **Loop** from $i = 0$:

- **For** each state $s \in S$:

$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$

- **Until** $\max_s |V_{i+1}(s) - V_i(s)| < \textit{small threshold}$

Convergence of Value Iteration

- The Bellman update is a **contraction mapping**
- Fact 1: Bellman update does not change optimal values V^* (*fixed point*)

$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Fact 2: The *max norm* $\|V_i - V^*\| = \max_s |V_i(s) - V^*(s)|$ between any value function V_i and V^* satisfies

$$\|V_{i+1} - V^*\| \leq \gamma \|V_i - V^*\|$$

- Each update shrinks “error” in V by factor of γ !

Rate of Convergence


- Value iteration converges *exponentially* fast
- Recall that a state's value is bounded by $\frac{|r_{\max}|}{1-\gamma}$
- After k passes of value iteration, error is bounded by $\gamma^k \frac{|r_{\max}|}{1-\gamma}$
- The smaller the γ , the faster that the error shrinks
- Tradeoff: Decisions become more myopic, future rewards less appealing

Values to Policy

- The goal of value iteration is to eventually extract an optimal policy
- We already know how to find π^* given V^* :

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Observation: We don't *require* optimal values
- Only relative values matter
- Policy may converge long before values do



Policy Iteration

- Idea: A policy can be computed at any point during value iteration
- We can improve on policy *directly*, leading to better values, leading to a better policy, and so on...
- Initialize $\pi_1(s)$ arbitrarily, $V^{\pi_0}(s) \leftarrow 0$ for all states s
- **Loop** from $i = 1$:
 - **Policy evaluation:** Compute V^{π_i} using $V^{\pi_{i-1}}$ as initial values
 - **Policy improvement:** From V^{π_i} , find new policy π_{i+1}
- **Until** $\pi_{i+1} = \pi_i$

Policy Evaluation


$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- Review: System of $|S|$ linear equations in $|S|$ unknowns

- Alternatively, take an iterative approach:

$$V_{k+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$

- This is just value iteration without max!
- Can be faster if initialized with values of similar policy




Policy Improvement

- Given values for a fixed policy, how can we improve it?
- Consider taking “greediest” action at each state:

$$\pi_{i+1}(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

- If π_i already optimal, then $V^{\pi_i} = V^*$ and $\pi_{i+1} = \pi_i$
- Otherwise, V^{π_i} can be moved closer to V^* by changing some actions
- Updating policy using argmax analogous to updating values using max



Example: Mini-Gridworld

- States: A, B, C ; actions: L, R
- Rewards received when entering state
- Transitions: $\Pr(\text{intended direction}) = 0.8, \Pr(\text{opposite direction}) = 0.2$

+3	-2	+1
----	----	----


- Suppose we initialize $(\pi_0(A), \pi_0(B), \pi_0(C)) = (R, R, R)$
- Evaluate** policy (either solve linear eqs or iterate Bellman-style):

$$(V_0(A), V_0(B), V_0(C)) = (-0.333, 1.75, 0.958) \quad \gamma = 0.5$$


- Improve** policy:

$$\begin{pmatrix} \pi_1(A) \\ \pi_1(B) \\ \pi_1(C) \end{pmatrix} = \begin{pmatrix} \operatorname{argmax}[0.8(3 + 0.5V_0(A)) + 0.2(-2 + 0.5V_0(B)), -0.333] \\ \operatorname{argmax}[(0.8(3 + 0.5V_0(A)) + 0.2(1 + 0.5V_0(C))), 1.75] \\ \operatorname{argmax}[(0.8(-2 + 0.5V_0(B)) + 0.2(1 + 0.5V_0(C))), 0.958] \end{pmatrix} = \begin{pmatrix} L \\ L \\ R \end{pmatrix}$$

Value Iteration vs Policy Iteration



- Computes values only
- Keeps track of values only
- Each sweep consists of one iterative policy evaluation (sum) and policy improvement (max)



- Computes values and policy
- Keeps track of policy only
- Each sweep consists of many iterative policy evaluations (sum) and policy improvement (argmax)

Algorithm Complexity

- Each sweep of value iteration takes $O(|S|^2|A|)$ time

$$V_{i+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i(s')]$$

- Each sweep of policy iteration takes $O(|S|^3 + |S|^2|A|)$ time


$$V_{k+1}^{\pi_i}(s) \leftarrow \sum_{s'} T(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

$$\pi_{i+1}(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^{\pi_i}(s')]$$

- In practice, complexity is also strongly dependent on problem at hand

Algorithm Complexity

- Value iteration: $O(|S|^2|A|)$
- Policy iteration: $O(|S|^3 + |S|^2|A|)$
- Value iteration: Number of sweeps depends on γ and error threshold
- Increases dramatically for high discount factor
- Policy iteration: Policy evaluation typically much more efficient than $O(|S|^3)$
- Fewer sweeps needed overall to converge



Summary

- Dynamic programming solves MDPs exactly by using recursive relationships among the state values
- Bellman updates push values and policies toward the optimal solution
- Value iteration: Compute and converge toward optimal values for all states, then extract policy
- Policy iteration: Alternate between evaluating a current policy and improving the policy